



Lic. Informatique, Univ Lille 1, 2010-11

■ Introduction aux bases de données relationnelles

- 6ème séance: **SQL**, niveau intermédiaire

■ Enseignante: C. Kuttler

■ **Biblio**: chapitre 4 de *Database Systems Concepts* de Silberschatz et al, McGraw-Hill (6ème édition, 2010)

- Ces transparents sont une adaptation de ceux disponibles sur le site du livre: **www.db-book.com**



Chapter 4: Intermediate SQL

- Null values in queries (Chap 3)
- Join Expressions
- Table creation (already seen on 10/09 Chap 3)
- Table modification and updates (Chap 3)
- Integrity Constraints
- SQL Data Types and Schemas
- Authorization
- Views
- Transactions



Null Values: the Tricky Details

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies **an unknown value or that a value does not exist**.
- The result of any **arithmetic** expression involving *null* is *null*
 - Example: $5 + \text{null}$ returns null
- The predicate **is null** can be used to check for null values.
 - Example: Find all instructors whose salary is null.

```
select name  
from instructor  
where salary is null
```



Three Valued Logic: true, false, unknown

- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*
- Any **comparison** with *null* returns *unknown*
 - Example: $5 < \text{null}$ or $\text{null} \diamond \text{null}$ or $\text{null} = \text{null}$
- Three-valued logic using the truth value *unknown*:
 - **OR**: $(\text{unknown} \text{ or } \text{true}) = \text{true}$,
 $(\text{unknown} \text{ or } \text{false}) = \text{unknown}$
 $(\text{unknown} \text{ or } \text{unknown}) = \text{unknown}$
 - **AND**: $(\text{true} \text{ and } \text{unknown}) = \text{unknown}$,
 $(\text{false} \text{ and } \text{unknown}) = \text{false}$,
 $(\text{unknown} \text{ and } \text{unknown}) = \text{unknown}$
 - **NOT**: $(\text{not } \text{unknown}) = \text{unknown}$
 - “*P* is unknown” evaluates to true if predicate *P* evaluates to *unknown*



Null Values and Aggregates



Null Values and Aggregates

- Total all salaries

```
select sum (salary)  
from instructor
```



Null Values and Aggregates

- Total all salaries

```
select sum (salary )  
from instructor
```

- Above statement ignores null amounts



Null Values and Aggregates

- Total all salaries

```
select sum (salary)  
from instructor
```

- Above statement ignores null amounts
- Result is *null* if there is no non-null amount



Null Values and Aggregates

- Total all salaries

```
select sum (salary)  
from instructor
```

- Above statement ignores null amounts
- Result is *null* if there is no non-null amount
- All aggregate operations except **count(*)** ignore tuples with null values on the aggregated attributes



Null Values and Aggregates

- Total all salaries

```
select sum (salary)  
from instructor
```

- Above statement ignores null amounts
- Result is *null* if there is no non-null amount
- All aggregate operations except **count(*)** ignore tuples with null values on the aggregated attributes
- What if collection has only null values?



Null Values and Aggregates

- Total all salaries

```
select sum (salary)  
from instructor
```

- Above statement ignores null amounts
- Result is *null* if there is no non-null amount
- All aggregate operations except **count(*)** ignore tuples with null values on the aggregated attributes
- What if collection has only null values?
 - count returns 0



Null Values and Aggregates

- Total all salaries

```
select sum (salary)  
from instructor
```

- Above statement ignores null amounts
- Result is *null* if there is no non-null amount
- All aggregate operations except **count(*)** ignore tuples with null values on the aggregated attributes
- What if collection has only null values?
 - count returns 0
 - all other aggregates return null



Chapter 4: Intermediate SQL

- Null values in queries (Chap 3)
- Join Expressions
- Table creation (already seen on 10/09 Chap 3)
- Table modification and updates (Chap 3)
- Integrity Constraints
- SQL Data Types and Schemas
- Authorization
- Views
- Transactions



Joined Relations

- **Join operations** take two relations and return as a result another relation.
- A join operation is a Cartesian product which requires that tuples in the two relations match (under some condition). It also specifies the attributes that are present in the result of the join.
- **Join condition** – defines which attributes in the two relations match
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

<i>Join types</i>	<i>Join Conditions</i>
inner join left outer join right outer join full outer join	natural on <predicate> using (A_1, A_1, \dots, A_n)



Natural Join

- Natural join matches tuples with the **same values for all common attributes**, and retains only one copy of each common column
- **select ***
from instructor natural join teaches;

ID	name	dept_name	salary	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	CS-347	1	Fall	2009
12121	Wu	Finance	90000	FIN-201	1	Spring	2010
15151	Mozart	Music	40000	MU-199	1	Spring	2010
22222	Einstein	Physics	95000	PHY-101	1	Fall	2009
32343	El Said	History	60000	HIS-351	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-101	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-319	1	Spring	2010
76766	Crick	Biology	72000	BIO-101	1	Summer	2009
76766	Crick	Biology	72000	BIO-301	1	Summer	2010



Join operations – More examples

■ Relation *course*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

■ Relation *prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

- Note: no prerequisite for CS-315, and no course Information for CS-437.



Inner join

- *course* **inner join** *prereq* **on**
course.course_id = prereq.course_id

Correction: *prereq_id*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

- **Equivalent to:**

select * from *course*, *prereq* **where** *course.course_id = prereq.course_id*

course **join** *prereq* **on** *course.course_id=prereq.course_id*



Outer Join

- Avoids loss of information.
- Computes the join
 - then adds tuples from one relation for which there is no match in the other relation to the result of the join.
 - Fills the empty attributes in the result with *null* values.



Left Outer Join

- *course* **natural left outer join** *prereq*
- *CS-315 doesn't have any prerequisite!*
 - *Corresponding line padded with null value*

Correction: **prereq_id**

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>





Right Outer Join

- *course* **natural right outer join** *prereq*
- course CS-347, occurring in the prerequisite table, isn't listed in the course table! Corresponding attributes are padded with null values in result.

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101



Full Outer Join

■ *course* **natural full outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101



Questions

- How to use outer joins on to detect anomalies on INFOTOUR?
 - Find the identifier of a director who's not listed in the PERSONNES table.
 - Check for other possible violations of what is shown as foreign key constraints in INFOTOUR's diagram.
- How can we obtain the same information, about tuples of one table for which there is no matching data in another table, without outer joins?
- Other useful queries on the university database with outer joins?



Chapter 4: Intermediate SQL

- Null values in queries (Chap 3)
- Join Expressions
- Table creation (already seen on 10/09)
- Table modification and updates (Chap 3)
- Integrity Constraints
- SQL Data Types and Schemas
- Authorization
- Views
- Transactions



Create Table Construct

- An SQL relation is defined using the **create table** command:

```
create table r (A1 D1, A2 D2, ..., An Dn,  
                (integrity-constraint1),  
                ...,  
                (integrity-constraintk))
```

- *r* is the name of the relation
- each *A*_{*i*} is an *attribute name* in the schema of relation *r*
- *D*_{*i*} is the *data type* of values in the domain of attribute *A*_{*i*}

- Example:

```
create table instructor (  
    ID           char(5),  
    name        varchar(20) not null,  
    dept_name   varchar(20),  
    salary      numeric(8,2))
```




Integrity Constraints in Create Table

- not null
- default value
- primary key (A_1, \dots, A_n)
- foreign key (A_1, \dots, A_n) references r

Example. Declare *branch_name* as the primary key for *branch*

```
create table instructor (  
    ID          char(5),  
    name       varchar(20) not null,  
    dept_name varchar(20),  
    salary     numeric(8,2),  
    primary key (ID),  
    foreign key (dept_name) references department))
```

primary key declaration on an attribute automatically ensures **not null**



And a Few More Relation Definitions

- **create table** *student* (
 ID **varchar(5) primary key**,
 name **varchar(20) not null**,
 dept_name **varchar(20)**,
 tot_cred **numeric(3,0)**,
 foreign key (*dept_name*) **references** *department*));

- **create table** *takes* (
 ID **varchar(5) primary key**,
 course_id **varchar(8)**,
 sec_id **varchar(8)**,
 semester **varchar(6)**,
 year **numeric(4,0)**,
 grade **varchar(2)**,
 foreign key (*ID*) **references** *student*,
 foreign key (*course_id, sec_id, semester, year*) **references** *section*);



Chapter 4: Intermediate SQL

- Null values in queries (Chap 3)
- Join Expressions
- Table creation (already seen on 10/09 Chap 3)
- Table modification and updates
- Integrity Constraints
- SQL Data Types and Schemas
- Authorization
- Views
- Transactions



Drop and Alter Table Constructs

- **drop table**
- **adding and removing attributes to existing table**
 - **alter table r add $A D$**
 - ▶ where A is the name of the attribute to be added to relation r and D is the domain of A .
 - ▶ All tuples in the relation are assigned *null* as the value for the new attribute.
 - **alter table r drop A**
 - ▶ where A is the name of an attribute of relation r
 - ▶ Dropping of attributes not supported by many databases.
- **adding constraints to existing table**
 - **alter table table-name add constraint**
 - System first ensures that the relation satisfies the constraint. If it does, the constraint is added to the relation. If not, the



Modification of the Database – Insertion

- Add a new tuple to *course*

insert into *course*

values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);

- or equivalently

insert into *course* (*course_id*, *title*, *dept_name*, *credits*)

values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);

- Add a new tuple to *student* with *tot_creds* set to null

insert into *student*

values ('3003', 'Green', 'Finance', *null*);



Modification of the Database – Insertion 2

- Insertion of tuples based on an existing table
- Add all instructors to the *student* relation with *tot_creds* set to 0

insert into *student*

select *ID, name, dept_name, 0*
from *instructor*

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation
- (otherwise queries like
insert into table1 select * from table1
would cause problems)



Modification of the Database – Updates

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others receive a 5% raise

- Write two **update** statements:

```
update instructor
  set salary = salary * 1.03
  where salary > 100000;
update instructor
  set salary = salary * 1.05
  where salary <= 100000;
```

- The order is important
- Someone earning just under 10000\$ may obtain two increases!
- Can be done better using the **case** statement (next slide)



Case Statement for Conditional Updates

- Same query as before but with case statement

```
update instructor  
  set salary = case  
    when salary <= 100000 then salary * 1.05  
    else salary * 1.03  
  end
```




Updates with Scalar Subqueries

- **Scalar subquery: returns a value (and not: a table)**
- Recompute and update *tot_creds* value for all students

update *student S*

```
  set tot_cred = ( select sum(credits)  
                  from takes natural join course  
                  where S.ID= takes.ID and  
                      takes.grade  $\neq$  'F' and  
                      takes.grade is not null);
```

- Sets *tot_creds* to null for students who have not taken any course
- Instead of **sum**(*credits*), use:

```
  case  
    when sum(credits) is not null then sum(credits)  
    else 0  
  end
```



Modification of the Database – Deletion

- Delete all instructors

delete from *instructor*

- Delete all instructors from the Finance department

delete from *instructor*
where *dept_name* = 'Finance';

- Delete all tuples in the *instructor* relation, for those instructors associated with a department located in the Watson building.

delete from *instructor*
where *dept name* in (**select** *dept name*
 from *department*
 where *building* = 'Watson');



Example Query

- Delete all instructors whose salary is less than the average salary of instructors



Example Query

- Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor  
where salary < (select avg (salary) from instructor);
```



Example Query

- Delete all instructors whose salary is less than the average salary of instructors

delete from *instructor*
where *salary* < (**select avg** (*salary*) **from** *instructor*);

- Problem: as we delete tuples from deposit, the average salary changes
- Solution used in SQL:
 1. First, compute **avg** salary and find all tuples to delete
 2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)



Chapter 4: Intermediate SQL

- Null values in queries (Chap 3)
- Join Expressions
- Table creation (already seen on 10/09 Chap 3)
- Table modification and updates (Chap 3)
- **Integrity Constraints**
- SQL Data Types and Schemas
- Authorization
- Views
- Transactions



Integrity Constraints

- Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency.
 - A checking account must have a balance greater than \$10,000.00.
 - A salary of a bank employee must be at least \$4.00 an hour.
 - A customer must have a (non-null) phone number.



Constraints on a Single Relation

- **not null**
- **default value**
- **primary key**
- **unique**
- **check (P)**, where P is a predicate



Not Null and Unique Constraints

■ not null

- Declare *name* and *budget* to be **not null**

name **varchar(20) not null**

budget **numeric(12,2) not null**

■ unique (A_1, A_2, \dots, A_m)

- The unique specification states that the attributes A_1, A_2, \dots, A_m form a candidate key.
- Candidate keys are permitted to be null (in contrast to primary keys).



The check clause

■ **check** (P)

where P is a predicate

Example: ensure that semester is one of fall, winter, spring or summer:

```
create table section (  
    course_id varchar (8),  
    sec_id varchar (8),  
    semester varchar (6),  
    year numeric (4,0),  
    building varchar (15),  
    room_number varchar (7),  
    time slot id varchar (4),  
    primary key (course_id, sec_id, semester, year),  
    check (semester in ('Fall', 'Winter', 'Spring', 'Summer'))  
);
```



Referential Integrity

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
 - Example: If “Biology” is a department name appearing in one of the tuples in the *instructor* relation, then there exists a tuple in the *department* relation for “Biology”.
- Let A be a set of attributes. Let R and S be two relations that contain attributes A and where A is the primary key of S. A is said to be a **foreign key** of R if for any values of A appearing in R these values also appear in S.



Cascading Actions in Referential Integrity

- **create table** *course* (
 course_id **char**(5) **primary key**,
 title **varchar**(20),
 dept_name **varchar**(20) **references** *department*
)

- **create table** *course* (
 ...
 dept_name **varchar**(20),
 foreign key (*dept_name*) **references** *department*
 on delete cascade
 on update cascade,
 ...
)

- alternative actions to cascade: **set null, set default**



Integrity Constraint Violation During Transactions

■ E.g.,

```
create table person (  
    ID char(10),  
    name char(40),  
    mother char(10),  
    father char(10),  
    primary key ID,  
    foreign key father references person,  
    foreign key mother references person)
```

■ How to insert a tuple?

■ What if *mother* or *father* is declared not null?

- **constraint** *father_ref* **foreign key** *father* **references** *person*,
 constraint *mother_ref* **foreign key** *mother* **references** *person*)
- **set constraints** *father_ref*, *mother_ref* **deferred**



Complex Check Clauses

- **check** (*time_slot_id* in (**select** *time_slot_id* from *time_slot*))
 - why not use a foreign key here?
- Every section has at least one instructor teaching the section.
 - how to write this?
- Unfortunately: subquery in check clause not supported by pretty much any database
 - Alternative: triggers (later)
- **create assertion** <assertion-name> **check** <predicate>;
 - Also not supported by anyone



Chapter 4: Intermediate SQL

- Discussed today:
- Null values in queries (Chap 3)
- Join Expressions
- Table creation (already seen on 10/09 Chap 3)
- Table modification and updates (Chap 3)
- **Integrity Constraints**